

Analyzing Topic Prevalence in Popular Hacker News Stories

Jake J. Dalli

Department of Artificial Intelligence

Faculty of ICT

University of Malta

jake.dalli.10@um.edu.mt

Abstract—The online link aggregator Hacker News has emerged as one of the prime rallying points for startup founders, venture capitalists and technology workers; grasping the attention of this cohort of internet users carries intrinsic value. In this paper, we should how certain topics more likely to make it to the front page and that specific topics are definitely over-represented within the top post. We frame the problem as a topic mining challenge, applying a bayesian generative statistical model, Latent Dirichlet Allocation (LDA), and subsequently analyzing it to draw our conclusions.

Index Terms—Statistics, Machine Learning, Data Mining, News, Aggregator, Topic Mining, Latent Dirichlet Allocation, Natural Language Processing

I. INTRODUCTION

Founded by venture capitalist Paul Graham in 2007, Hacker News (HN) is a news aggregator which lies at the heart of the startup renaissance [1]. While initially tailored towards startup founders, HN has come to represent more than just startup information, with participants urged to post “*anything that gratifies one’s intellectual curiosity*.” In this paper, we perform an analysis of the front page with particular attention to the temporal and topical aspects of the data at hand [2].

Our results indicate that certain topics related to corporate revenue growth, such as making money, corporate transactions and recruitment; and topics related to programming niches such as Rust and Ruby programming, concurrency and parallel programming and interpreted languages tend to make it to the front page more often. We also conclude that posts related to web browsers tend to be over-represented within the top ranked story.

II. APPROACH

We aim to investigate the composition of the front-page of the aggregator, looking to understand whether any topics are particularly prevalent, and what other

attributes a popular story might possess. Thus, we ask the following questions:

- 1) Is it easier to make it to the front page of Hacker News if you publish an article on a specific day or at a specific time?
- 2) Are there particular topics which make it more often to the front page?
- 3) Are some topics over-represented within the top post?

The task of detecting topics lends itself towards the field of generative models, due to the fact that we are looking to generate a number of outcomes (in this case topics) based on our input observations [3]. Moreover, we also consider that our problem domain may result in noisy or incomplete data. Apart from news articles, HN users also post links to videos, unconventional web pages and web applications (this is particularly prevalent for stories tagged *Show HN*) amongst others.

Our analysis takes two different forms; we first perform a **descriptive analysis** of the dataset, investigating the statistical relationship between points and comments, and the relationship between the date and time that the story was published in relation to the rank of the story; secondly, we perform **topic detection** by building a *bayesian generative statistical model* to analyze the topics detected [4].

With these goals in mind, we make two key assumptions by analyzing the hackernews datasource:

- We shall limit our analysis to the top 30 HN stories for a given day, ranked by the number of points the article received. This is distinctively different to analyzing the articles which appear on the front page. The HN front-page is composed both of top ranked articles, and new articles. Typically new articles quickly make their way out of the front page if they do not receive enough points to maintain their presence. Thus one key assumption is that the top 30 articles ranked by points have appeared on

the front page. This assumption allows us to draw more accurate results by eliminating noise from the dataset.

- We limit our approach to detecting topics within text-based stories, purposefully excluding or knowingly disregarding stories which link to files (such as PDF documents), web applications or videos. Our approach only considers text articles.

A. Latent Dirichlet Allocation

We approach the task of topic detection by building a generative model using *Latent Dirichlet Allocation*, which represents documents as mixtures of latent topics characterized by term distributions [5]. The model assumes the following process for each document d in a corpus C :

- 1) Choose the number of terms N assigned to a document according to a Poisson distribution ϕ
- 2) Choose a topic mixture θ according to Dirichlet distribution α over a fixed number of K topics.
- 3) For each N :
 - a) First pick a topic Z_n based on a multinomial distribution sampled from θ .
 - b) Based on Z_n generate a term w_n from $p(w_n|Z_n, \beta)$, computing a multinomial probability conditioned on z_n

To illustrate this, consider a corpus containing documents about company transactions. A document headlined "*Salesforce acquires Tableau for \$15.8B*" may yield the terms *Salesforce*, *acquire*, *Tableau* ($N = 3$). We may have two topics ($K = 3$) about *mergers*, *acquisitions* and *restructurings*. LDA would generate the three terms with their relative probability in relation to the topics; so for example the term *Salesforce* would be associated to the three topics *mergers*, *acquisitions* and *restructurings* with probabilities $P_m = 0.3$, $P_a = 0.9$ and $P_n = 0.5$ respectively.

Given enough documents which outline several acquisitions by *Salesforce*, we may discover a topic which specifically relates to company transactions by the company.

The LDA process assumes prior knowledge of the number of topics, K which are to be detected. This can prove challenging in applications where the number of topics is not known a priori. Thus the LDA approach is tightly coupled with a parameter estimation problem for deducing the ideal number of topics.

Topic mixture can be evaluated by computing the *Maximum Log Likelihood* of the given data [6]. This is done by holding out a test set of unseen documents D_t ,

to be evaluated against a topic matrix Υ , with a hyper-parameter α . Thus the Log-Likelihood can be computed as follows:

$$L(W) = \log P(d|\Upsilon, \alpha) = \sum_t \log P(d_t|\Upsilon, \alpha) \quad (1)$$

One can thus measure the *perplexity* of a topic model:

$$\text{perplexity}(w_{n..t}) = \left\{ -\frac{L(W)}{\text{TotalTokens}} \right\} \quad (2)$$

There are some observations to be made when evaluating topic mixtures. Firstly, the concept of a topic is highly coupled to human reasoning which is not quantitatively justifiable, thus perplexity and model quality may not be correlated [7]. Secondly, the *Log Likelihood* calculation is prone to local maxima, thus one must adjust additional hyper parameters to skip or discard certain iterations.

B. Data Collection and Storage

The Hacker News interface is updated in real-time with no pages showing the history of story submission. Each story is accompanied by a story headline, story publish date, number of points and comments, story link, thread link and source domain.

We perform a one-time web page extract from the meta-aggregator HckrNews [8], a website which displays an archive of the top voted HN stories by date and time. Our extract consists of the top 50% of stories on the HN board for six months ranging from *1st December 2018* till *31st May 2019*. In turn, we design a pipeline of R scripts which extract the following data sources:

- 1) **List of Stories:** Parse the HTML page for stories, output a clean dataset with the aforementioned fields in comma delimited format (CSV). Filter out only the top 30 stories for each day by points.
- 2) **Source Pages:** For every story, visit the story link and download the web page in raw HTML format. This results in a repository of HTML pages accessed via the source link.
- 3) **Source Datasets:** Two datasets containing the stories extracted; *hackernews board*, containing only data shown on the bulletin board for the top 30 hackernews stories between 1st December 2018 and 31st May 2019; and *hackernews board content* which additionally contains the corpora (source pages and story headline keywords) which we shall use for topic processing and analysis. We explain the steps for cleaning and aggregating such data within the subsequent section.

Our results are stored in a dataset called *final-annotated-topic-document-matrix.CSV*, which contains

every story and its probability of association to each one of the 114 detected topics. It also contains the topic meta-data for the topics with the top three probabilities.

C. Data Preparation and Cleaning

Our preparation and cleaning processes perform three functions; firstly we perform some superficial cleaning and preparation of the *story list dataset*, by removing sub-strings related to website aesthetics. Secondly, we carry out information extraction on the source page HTML documents, to extract clean text from the HTML markup. Lastly, we perform a number of *Natural Language Processing* preparation techniques on the extracted repositories to reduce the volume of the data, and refine the textual representation in order to render better results.

The latter process proves to be nuanced, since it requires taking some strategic decisions to reduce the volume of data, and to remove linguistic features which may not be good ontological indicators. We assume that a document's ontological composition is adequately represented by its *noun phrases* and *verb phrases*. Another assumption is that the first 800 words of a document represents the ontological composition of a document, this is done to reduce the data volume. We also strip text corpora of punctuation and stop-words. Finally, we perform word stemming on the dataset.

We make use of the *BoilerpipeR R library* [9] to detect and remove template code from the HTML pages, which provides a model specifically trained for extracting articles from news website templates; furthermore, we use the *Text Mining and Open NLP R libraries* to clean the text data [10] [11].

In summary, we perform the following processes on the data:

- 1) Extract the story name, points, comments, source domain, source link, thread link, rank and publish date from the HckrNews [8] archive:
 - Remove the source domain from the story name.
 - Remove any trailing white-spaces
 - Convert the publish date to a readable date format.
 - Save the dataset as *hackernewsboard.csv*
- 2) For every story, visit the source domain and download the HTML page to the source article repository.
- 3) For each story, read the relevant documents from the source page repository. Extract nouns and verbs from the story text and story headline and save to separate fields within the data frame. Save this dataset as *hackernewscontent.csv*.

D. Methodology

Our analysis takes two forms. Within our descriptive analysis, we aim to answer the following questions:

- 1) *Is there a correlation between the number of points and the number of comments a story receives? Are some data sources more successful than others?*
- 2) *Does the time of day and day of week affect the ranking a story receives?*

The second part of our analysis involves training a Latent Dirichlet Allocation Model trained on the extracted data. Our model is trained with a corpus containing of documents representing the *Noun Phrases* and *Verb Phrases* within the story headline, and the first 800 words of the source page.

We do this by constructing a Document Term Matrix (TF-IDF) and training the model using *Gibbs Sampling* to find the conditional probability distribution of a word's topic assignment, modifying θ to achieve the maximum log likelihood. However, this method assumes that we already know the number of topics we need to detect.

Thus, we must perform a *parameter selection process*, repeatedly running the model with different values of K , and evaluating the topic mixture by calculating the *Maximum Log Likelihood*. Each model is run using 700 iterations, discarding the first 200 iterations to increase accuracy. The results for this process are reported within our analysis.

When the ideal k is selected, we train our model once again to preform our analysis, answering the following questions:

- *Are there particular topics which make it more often to the front page?* We build a document-topic probability correlation covariance matrix representing the correlation of two topics on our document. Next, we cluster the topics based on their covariance using hierarchical clustering. The clusters detected are indicative of groups of topics appearing more frequently on the front page.
- *Are some topics over-represented within the top post?* We create a bucket for the rank attribute, marked *true* if the article is top ranked, and *false* if the article has any other rank. Next, we perform a Chi Square Test [12] to investigate the correlation between a top ranked article and the topic. To ascertain our results, we also perform a One-Way ANOVA test [13] on the ranking as an ordinal value.

The primary method for topic detection used is Latent Dirichlet Allocation, as provided by the *Topic Models R library* [14]; We visualize our data using the *GGPlot R library* [15], using a theme derived from that produced by The Economist magazine.

III. DESCRIPTIVE ANALYSIS

We begin our analysis by plotting stories and sources. The plots are densely populated at the zero axis, becoming less densely populated as they increase in points and comments. This is indicative of the fact that higher story rankings are more difficult to achieve.




Fig. 1. Stories plotted by Comments and Points within density bins

Perhaps unsurprisingly, we find that the relationship between the number of points assigned to an article, and the number of comments of an article is moderately correlated with a Pearson's R of 0.6 [16].

The plot is gently skewed towards two directions - higher rated posts tend to receive fewer comments, while highly commented posts tend to receive less points. While strongly correlated, we observe that the regression line may not be as steep as expected because not every user who assigns an up-vote may comment and not every user who comments may assign an up-vote.




Fig. 2. Stories plotted by Comments and Points grouped by Ranking

Laying motive for answering our second question regarding over-representation within the top post, we

observe that articles with lower points tend to receive more comments than articles with higher points.

To further illustrate this point, we plot the regression line for articles ranked within the top 5, and all other articles, showing a gentle reduction between point-comment correlation. In fact, for lower ranked articles, the correlation is weakly correlated.

We also note that the correlation of the *other* group is larger than that for articles ranked four and five, we attribute this to the size of the group. Given a plot which is further segmented, we would continue to see a reduction in the correlation until it flattens out.

Our observations are further corroborated by plotting the density curve for points and comments. Not only do stories tend to receive less comments than points, but we also observe that there is a larger variation in the number of comments an article receives, as opposed to the number of points.

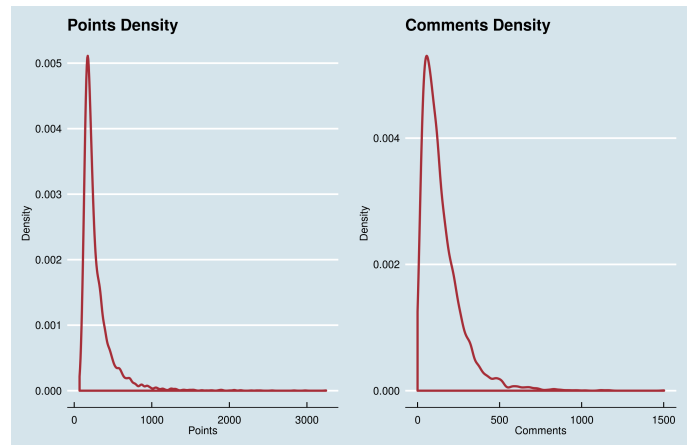


Fig. 3. Density plot for points and comments

We extend our analysis of points and comments correlation to sources, noting that the correlation is a strong correlation between points and comments for specific sources. The most widely commented and ranked sources are perhaps unsurprising, including popular news outlets such as the New York Times, Bloomberg, Techcrunch and social websites such as Github, Twitter and Medium. We note that the correlation is clearer in this case primarily due to the fact that we are considering data sources which occur more frequently within the top hacker news rankings. A news outlet of the size and magnitude such as that of the New York Times, publishes more frequently than a lone developer's blog, thus it stands to reason that it ranks high within our analysis.

However, the same plot for the *median* number of points gives a very different picture. This plot is more representative of what can be considered to be quality Hacker News data sources, since it is not skewed by the




Fig. 4. Source Domains plotted by Sum of Comments and Points

frequency of stories. In fact, the correlation is similar to that presented in Figure 1. In this plot, we may observe




Fig. 5. Source Domains plotted by Median of Comments and Points

data sources with very few stories published to the board, which have been a major success serving as *one hit wonders* of sorts. While GitHub retains its status as a top quality source on hacker news, the top voted source within our dataset is *deepsouthventures.com* which has a very successful story related to bootstrapping an Onion sales startup. Similarly, it is followed by *smartcar.com*, a startup whose product was cloned by a heavily funded competitor.

Given the variety of data sources, we can qualitatively conclude that Hacker News serves its mission in providing articles which stimulate curiosity over simple corporate content.

We can conclude that content rules over quantity and power of association, however we also observe that the time of day is also indicative of the number of points a story receives.




Fig. 6. Heat-map of the median number of HN points by Month and Day of the Week

Next, we analyze stories based on the date and time being published. We note that a story may not necessarily be up-voted at the time of publishing, however stories do need to receive a number of up-votes in order to maintain their status on the front page; subsequently, the points on a story tend to compound as the hours pass. Thus, analyzing the date and time is necessary to understanding what constitutes a top post.




Fig. 7. Points by Hour

The heat-map shows that HN articles published on weekdays are more successful than stories published on weekends. This can be attributed to the fact that since it is a data source which caters for business-related news, users frequent the site more often during business hours. Alternatively, one can hypothesize that HN is considered to be a healthier workplace distraction than other social media websites, however this hypothesis is difficult to substantiate.

To investigate this phenomenon, we investigate the number of points by the hour of the day. A visual

inspection of the histogram of points shows that articles published later on in the day tend to be more successful than those published in the morning. Stories published between 2pm and 7pm tend to receive the most points.

The top posted articles tend to be posted after lunch time or during the night, possibly indicating that HN does indeed serve as a distraction, either from the workplace or from the bedtime. Finally, we extract a word cloud of terms from the story headlines, giving us an indication of what topics to expect within the topic analysis section; in particular, it drives our decision to limit our topic detection to texts containing noun phrases and verb phrases.

IV. TOPIC DETECTION

We train our LDA model using 5,939 documents containing 7,947 terms, observing that the dataset is highly sparse.

Document Term Matrix	
Documents	5,939
Terms	7,947
Sparsity	99%

TABLE I
DOCUMENT TERM MATRIX FOR LDA TRAINING

Next, we perform the *parameter selection process* by repeatedly running the model for K values one through 120, at 700 iterations with a burn-in rate of 200. After each run, we calculate the Maximum Log-Likelihood (LL) of the model to evaluate the topic mixture. We cease our testing when the Maximum LL begins to taper off, seemingly converging. Figure 9 shows the Maximum LL plotted against K , showing the ideal K to be 114 topics.

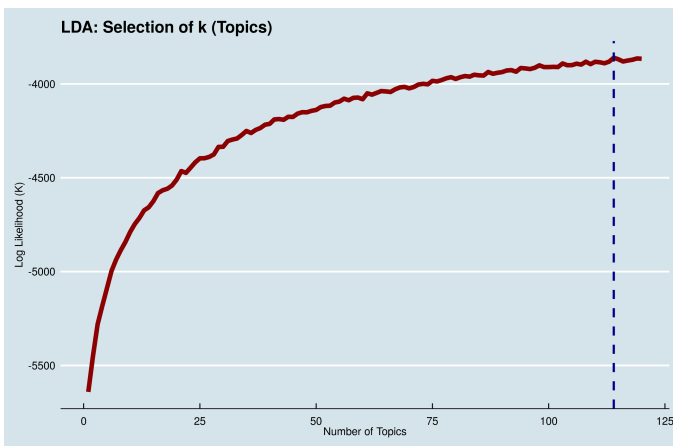


Fig. 8. K Selection Process

The LDA model outputs a series of topics, identified by their keywords. Each topic was evaluated, and annotated with a brief human-readable description in order to

facilitate analysis. Table II shows a number of examples. Most of the topics are clearly discernable as pertaining to a particular topic, however 18 of the 114 topics were identified as *noisy topics*, which did not have a single discernable topic.

Topic Id	Keywords	Human-Readable Annotation
1	point view channel youtube content	Youtube
2	flight plane air boe pilot	Boeing 737 & Aviation
3	countri india germani europ world	Geopolitics
4	phone devic call android smartphon	Mobile Devices & Apps
5	life friend mind world advic	Self Improvement
6	rust webassembl red swift runtim	Rust Lang.
7	materi glass plastic wast wood	Environmental Concerns
8	earth moon land planet star	Space Exploration
9	compani fund investor capit founder	VC, IPOs & Money
10	editor edit guid emac studio	IDEs
11	money peopl dollar million thousand	Making Money
12	peopl decis lot problem fact	Quitting Jobs & Toxic SV Culture
13	ve lot ll note	Noisy Topic 1
14	perform intel cpu processor core	CPUs, Performance & Advances
15	market busi industri compani sale	Corporate Revenue & Transactions

TABLE II
EXAMPLES OF HUMAN-READABLE ANNOTATIONS ASSIGNED TO LDA TOPICS

The topics detected span a wide range of subjects. Our qualitative annotation exercise reveals that Hacker News is a diverse bulletin board, referencing stories spanning from software development, consumer electronics, psychology, self-help, business and the environment. Some topics treat programming languages such as Rust and Ruby, while others deal with programming niches such as Numerical Computing, Machine Learning and Concurrency. Moreover, we observe a large number of topics relating to self-help, psychology, geo-politics and the environment.

Similar to our descriptive analysis, we plot topics by the median points and comments, showing us a visual indication of the most successful topics.




Fig. 9. Detected Topics by Comments and Points

A cursory view of the plot shows that the top topics are *Hiring and Careers*, encompassing stories related to the

workplace and the hiring process; *Making Money*, related to different businesses models and their revenues; *Web Protocols and APIs*, relating to web standards and news related to updates; *Wikileaks and Law*, pertaining to articles about *Julian Assange* and court orders; and *Apple*, the multinational technology company. Such topics are indicative of the fact that Hacker News brings together startup founders and software developers.

Annotation	Story Headline	Points
Hiring & Careers	Absolute truths I unlearned as junior developer (monicalent.com)	1478
Hiring & Careers	I interviewed at six top companies in Silicon Valley in six days (blog.usejournal.com)	970
Hiring & Careers	When hiring senior engineers, youre not buying, youre selling (hiringengineersbook.com)	941
Making Money	How to Be Successful (blog.samaltman.com)	836
Making Money	Firefox desktop market share now below 9% (netmarketshare.com)	821
Making Money	Open Source Doesnt Make Money Because It Isnt Designed to Make Money (www.ianbicking.org)	716
Web Protocols & APIs	Remote Code Execution on Most Dell Computers (d4stiny.github.io)	870
Web Protocols & APIs	HTTP headers for the responsible developer (www.twilio.com)	866
Web Protocols & APIs	HTTP/3 explained (http3-explained.haxx.se)	708
Wikileaks & Law	Julian Assange arrested in London (www.bbc.co.uk)	2354
Wikileaks & Law	U.S. Supreme Court Puts Limits on Police Power to Seize Private Property (www.nytimes.com)	1411
Wikileaks & Law	If Software Is Funded from a Public Source, Its Code Should Be Open Source (www.linuxjournal.com)	1131
Apple Corp.	Spotify to Apple: Time to Play Fair (www.timetoplayfair.com)	1876
Apple Corp.	FaceTime bug lets you hear audio of person you are calling before they pick up (9to5mac.com)	1531
Apple Corp.	Apple Sign In (techerunch.com)	1137

TABLE III

TOP STORIES WHICH FORM PART OF THE TOP VOTED TOPICS

1) *Topic Prevalence within the Front Page*: Following a visual analysis of the top topics, we aim to answer the question: "Are there particular topics which make it more often to the front page?" On one hand, one may consider topics in isolation, as per our prior analysis, however this does not account for the fact that topics are overlapping and inter-related. Moreover, the concept of a topic does not occur in isolation, on the contrary, a popular topic is bound to have other related popular topics. Thus we aim to detect clusters of topics, which occur more frequently together.

We do this by computing a document-topic probability covariance matrix, which displays the probability of two topics occurring together. Figure 11 shows the covariance matrix containing the distribution over log probabilities between topics, clustered hierarchically.

Interpreting the covariance is not a trivial exercise, however one can see that there are a number of topics which are tightly related from the density of colours. We use this visualization to detect groups of topics as clusters. Table IV shows a table of topics which were extracted from our interpretation.

One can observe that *Cluster 1* is heavily related to corporate developments, including the topics *Making Money*, *Corporate Revenue & Transactions*, *Corporate Tax* and some other generic topics relating to geopolitics




Fig. 10. Covariance Matrix of Topics

Topic Id	Topic Annotation	Cluster Name
11	Making Money	Cluster 1
15	Corporate Revenue & Transactions	Cluster 1
37	Corporate Tax	Cluster 1
23	C-Level Execs, Recruitment & Layoffs	Cluster 1
3	Geopolitics	Cluster 1
17	Lawmaking, Policy & Trump	Cluster 1
43	Payment Companies & Investments	Cluster 1
9	VC, IPOs & Money	Cluster 1
50	Government Surveillance	Cluster 1
30	Demographic Statistics	Cluster 1
51	Noisy Topic G	Cluster 1
48	Housing & Accom	Cluster 2
22	Japanese Tech.	Cluster 2
35	Problems & The Past	Cluster 2
21	Noisy Topic B	Cluster 2
20	Movies & Entertainment	Cluster 2
12	Quitting Jobs & Toxic SV Culture	Cluster 2
27	Terrorism & General Crime	Cluster 2
47	Math, Simulations & Trading	Cluster 3
53	Functional Languages	Cluster 3
44	Noisy Topic E	Cluster 3
6	Rust Lang.	Cluster 3
31	Noisy Topic C	Cluster 3
40	Ruby Lang.	Cluster 3
33	Interpreted Prog. Languages	Cluster 3
57	Concurrency & Parallel Programming	Cluster 3
26	Fast & Efficient Programming	Cluster 3
77	Nature, Animals & Insects	Cluster 4
113	Family Issues & Happiness	Cluster 4
94	Chess & Games & Board Games	Cluster 4
85	Noisy Topic K	Cluster 4
110	Noisy Topic Q	Cluster 5
89	Huawei & China	Cluster 5
58	Wikileaks & Law	Cluster 5
71	Corp. Lawsuits	Cluster 5

TABLE IV

CLUSTERS EXTRACTED FROM COVARIANCE MATRIX

and lawmaking. The second cluster, *Cluster 2*, contains topics which are related to lifestyle, including *Housing & Accommodation*, *Toxic Workplace Culture* and *Movies*.

Cluster 3 relates to software development niches, covering specific programming languages and techniques.

It is interesting to note that overlapping that topic 71, pertaining to corporate lawsuits, does not fall within *Cluster 1*. Moreover, we observe that our clustering method has detected many Noisy Topics - this can be attributed to the fact that the noisy topics do not focus on a single subject matter, and are likely to associated with several different documents.

Similar to our initial analysis, we plot these clusters by points and comments, in order to gain a better understanding of how the clusters rank against each other. We observe that *Cluster 1*, relating to corporate culture, generates the most activity both in terms of points and comments. However, the distribution is sparse. On the other hand, *Cluster 5*, which relates to censorship and lawsuits, is similarly successful but less sparse.

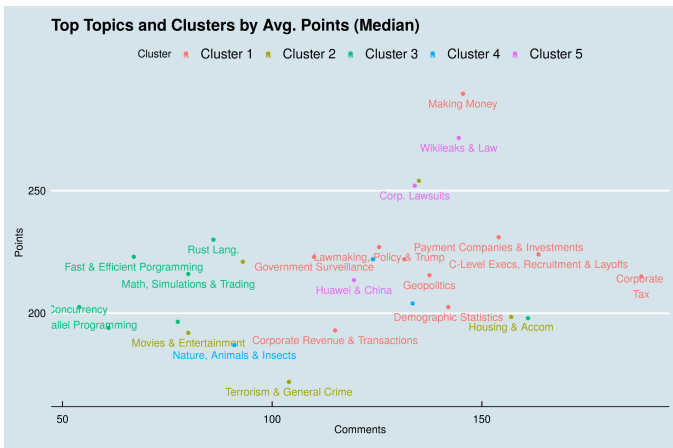


Fig. 11. Top Clustered Topics appearing on the Front Page

The cluster relating to niche programming concepts tends to be least discussed, but similarly successful to other clusters. It is also the most concentrated cluster within the plot. We also observe that clusters with more controversial topics, such as layoffs and tax, are more discussed - this is indicative of the fact that divisive topics create confrontation, and thus more dialogue.

We conclude that while our analysis shows several topic groups which are most likely to make it to the front page, topics relating to corporate entities, generating revenue, lawmaking, Wikileaks and niche software development are most represented on the front page.

2) *Topic Over-representation within the top post*: The above observations hold true for the whole front page of Hacker News, having detected topics from stories ranked one through thirty. Next, we aim to evaluate whether certain topics are over-represented within the top post. We pose the null hypothesis:

H_0 : "Achieving the top rank from within the front page is unrelated to the story topic."

We first perform a Chi-Square test between two categorical values: the topic name, and the binned rank, which shows a one if a story was in the top post, and zero if the story was shown in any other rank. A p -value of 0.002634, which falls under the 0.05 critical value is obtained.

One must note that the probability obtained with this test is remarkably small, this is primarily due to the reason that the frequencies are relatively small - thus our test may be inconsistent. Thus we perform an additional statistical test, a One-Way ANOVA Test. Our One-Way ANOVA test between a nominal value, the story rank, and a categorical value, the topic. Table V shows the results of our One-Way ANOVA test, resulting in a p -value which also falls under the critical value of 0.05.

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
Topic	113	16588	146.80	1.996	2.97e-09
Residuals	5825	428417	73.55		
Signifc. codes	0.001	0.01	0.05	0.1	1

TABLE V
ONE-WAY ANOVA TEST BETWEEN TOPICS (CATEGORICAL) AND RANK (NOMINAL)

We reject the null hypothesis H_0 with a p -value of 0.002634, which falls under the 0.05 critical p -value. Therefore we conclude that there is strong evidence that indicates that the top rank is *highly* dependent on the topics present, rejecting the null hypothesis.

Thus we are able to conclude that **sometopics are indeed over-represented within the top post** - we now investigate which topics are represented within the top post.




Fig. 12. Top Topics by Median Points

Filtering the topics by the top rank, and selecting the top topics by median post shows a different set of top

topics than observed across the full ranking - the only topic retained is related to Wikileaks and Law. The top ranked topic of all time is *Theoretical Physics*, due to the heavily up-voted article on the first ever image of a black hole. Table VI shows the list of top topics within the top rank using this method, while figure 12 shows a scatter plot of these topics.

Topic	Story	Points
Theoretical Physics	Unveiling the first-ever image of a black hole [video] (www.youtube.com)	2150
Academia	UC terminates subscriptions with Elsevier in push for open access (www.universityofcalifornia.edu)	1706
BizDev	Start with a Website, Not a Mobile App (www.atrium.co)	1536
Journo, Media & Ads	No Thank You, Mr. Pecker (medium.com)	2400
Journo, Media & Ads	Bezos Investigation Says the Saudis Obtained His Private Data (www.thedailybeast.com)	602
Containers, Orch & Cloud	Please do not attempt to simplify this code (github.com)	1483
Malicious Company Behaviour	Google Tried to Patent My Work After a Job Interview (patentpandas.org)	1757
Malicious Company Behaviour	Apps intended for kids may not include third-party advertising or analytics (developer.apple.com)	1171
Noisy Topic R	I made a smart watch from scratch (m.imgur.com)	1431
Wikileaks & Law	Julian Assange arrested in London (www.bbc.co.uk)	2354
Wikileaks & Law	U.S. Supreme Court Puts Limits on Police Power to Seize Private Property (www.nytimes.com)	1411
Wikileaks & Law	Justice Department Is Preparing Antitrust Investigation of Google (www.wsj.com)	867
Web Browsers & Plugins	Switch from Chrome to Firefox (www.mozilla.org)	3246
Web Browsers & Plugins	A Conspiracy to Kill IE6 (blog.chriszacharias.com)	2141
Web Browsers & Plugins	Google to restrict modern ad blocking Chrome extensions to enterprise users (9to5google.com)	2057
Self Improvement	Ask HN: What books changed the way you think about almost everything?	1905
Self Improvement	A guide to difficult conversations (medium.dave-bailey.com)	1782
Self Improvement	Why isn't the internet more fun and weird? (jarredsummer.com)	1495

TABLE VI

TOP TEN TOPICS IN THE TOP POST BY MEDIAN POINTS




Fig. 13. Topic Occurance in Top Post

We note that the top topics within the top rank tend to have a negative spin to them, relating to scandals, malicious company behavior and restrictions. However, one must note that this selection is highly prone to outliers. In fact, some of the top topics shown only have a one or two posts. This is perhaps due to the nature of the top post, thus we instead plot the *frequency* of topic occurrence within the top post.

Figure 13 and table VII show us that *Web Browsers & Plugins* is most likely to occur in the top rank, with a significantly higher probability than other topics.

In this analysis, we showed that some topics are indeed over-represented within the top rank, giving examples as to what these topics cover. Thus, we conclude that **some topics occur more often than others on the front page** and **some topics are over-represented within the top rank**. We summarize and elaborate on these key-findings in our conclusion.

Topic	Frequency	Probability
Web Browsers & Plugins	13	0.0660
Tool UI, Automation & Devops	8	0.0406
Github & New OSS	7	0.0355
Hiring & Careers	6	0.0305
Math & CS	6	0.0305
Noisy Topic D	6	0.0305
Self Improvement	6	0.0305
Software Usability & User Interfaces	6	0.0305
Apple Corp.	5	0.0254
Facebook Privacy Concerns	5	0.0254

TABLE VII

TOP FIVE TOPICS WITHIN THE TOP RANK BY PROBABILITY

V. CONCLUSION

Our analysis concludes that some topics are more prevalent than others on the Hacker News front page, and within the front page, certain topics are over-represented within the top post. Our key observations regarding the Hacker News front page are as follows:

- 1) The points a story receives and the comments a story receives are positively correlated, with a moderately to strong correlation and Pearson's R of 0.6.
- 2) Stories posted between 2pm and 7pm, and stories posted on weekdays tend to rank higher than others.
- 3) Certain topics tend to be over-represented on the front page. These include topics related to companies and revenue growth, such as making money,

corporate transactions, corporate tax, recruitment and policy; topics relating to programming niches such as Rust programming, Ruby programming, concurrency and parallel programming and interpreted languages; and topics relating to lifestyle such as housing, movies and entertainment, job mobility and workplace culture.

- 4) We reject the null hypothesis H_0 , which holds that achieving the top rank from within the front page is unrelated to the story topic, with a p-value of 0.002634, showing that there is strong evidence that indicates that the top rank is highly dependent on the topics it represents. Stories related to Web Browsers, such as Mozilla Firefox and Google Chrome, are over-represented within the top post.

The findings presented are the result of a topic detection process using Latent Dirichlet Analysis. A total of 114 topics were detected front page data spanning from the data collected, these topics were manually annotated and evaluated. Our evaluation process concluded that 18 of the 114 topics (15.8%) were identified as noisy, representing no discernable human-identifiable topic. Noisy topics mainly resulted from source pages which were not extracted correctly, due to pay-walls or scraping prevention mechanisms or pages which were not adequately cleaned. In other cases, some noisy topics contained conceptually similar articles spanning too generic a range to discern the common topic without reading the full source. The latter topics could not be justifiably marked as quality topics. Nonetheless, we conclude that the topic detection process was successful for the purposes of our analysis.

We attribute this success mainly to good decisions within our data cleaning process. In particular, the decision to represent articles by their *noun phrases* and *verb phrases* proved to be very effective.

Our analysis for investigating topic representation is based on the covariance matrix drawn from the topic model. We note that interpreting a large covariance matrix is a non-trivial task, and future work on the hierarchical clustering performed could render better quality topic clusters.

Lastly, our analysis for investigating over-representation within the top post resulted in a very low *p-value* for both *chi-squared* test and *one-way ANOVA*. We note that in the case of the *chi-squared* test, the p-value may not be entirely reliable mainly due to the low-frequencies present within the data [17].

We considered calculating correlation with a logistic regression, however we observed that this does not offer any significant advantages over a chi-squared test. Moreover, an alternative Fisher's Exact Test [18] [19] provided

similar results. Nonetheless, the fact that both our *chi-squared* test, which solely draws from the top rank, and the ANOVA test, which draws nominal values for all ranks, rendered similar results. This final observation increases our confidence in the results.

Further work may include investigating the seasonality of the detected topics, augmenting the LDA model with thread comments, and excluding corporate announcements. Such suggestions may serve to further corroborate our findings, and better capture the spirit of Hacker News.

REFERENCES

- [1] P. Graham, "Introducing startup news." <https://news.ycombinator.com/announcingnews.html>, Feb. 2007.
- [2] Ycombinator, "Hacker news guidelines." <https://news.ycombinator.com/newsguidelines.html>.
- [3] A. Y. Ng and M. I. Jordan, "On discriminative vs. generative classifiers: A comparison of logistic regression and naive bayes," in *Advances in neural information processing systems*, pp. 841–848, 2002.
- [4] M. S. J.K. Pritchard and P. Donnelly, "Inference of population structure using multilocus genotype data," *Genetics*, vol. 155, no. 2, pp. 945–959, 2000.
- [5] A. Y. N. David M Blei and M. I. Jordan, "Latent dirichlet allocation," *Journal of machine Learning research*, vol. 3, no. Jan, pp. 993–1022, 2003.
- [6] R. S. Hanna M Wallach, Iain Murray and D. Mimno, "Evaluation methods for topic models," in *Proceedings of the 26th annual international conference on machine learning*, pp. 1105–1112, ACM, 2009.
- [7] C. W. Jonathan Chang, Sean Gerrish, "Reading tea leaves: How humans interpret topic models," in *Advances in neural information processing systems*, pp. 288–296, 2009.
- [8] W. Larsen, "hckr news."
- [9] M. Annau, *Package boilerpipeR*. R Foundation for Statistical Computing.
- [10] I. Feinerer and K. Hornik, *tm: Text Mining Package*, 2018. R package version 0.7-6.
- [11] K. Hornik, "Cran - package opennlp."
- [12] M. L. McHugh, "The chi-square test of independence," *Biochemia medica: Biochemia medica*, vol. 23, no. 2, pp. 143–149, 2013.
- [13] R. M. Heiberger and E. Neuirth, *One-Way ANOVA*, pp. 165–191. New York, NY: Springer New York, 2009.
- [14] B. Grun, "Topic models [r package topicmodels version 0.2-8]."
- [15] L. H. Hadley Wickham, Winston Chang, "Create elegant data visualisations using the grammar of graphics [r package ggplot2 version 3.2.0]."
- [16] W. Kirch, ed., *Pearson's Correlation Coefficient*, pp. 1090–1091. Dordrecht: Springer Netherlands, 2008.
- [17] F. Yates, "Contingency tables involving small numbers and the chi 2 test," *Supplement to the Journal of the Royal Statistical Society*, vol. 1, no. 2, pp. 217–235, 1934.
- [18] R. A. Fisher, "On the interpretation of chi 2 from contingency tables, and the calculation of p," *Journal of the Royal Statistical Society*, vol. 85, no. 1, pp. 87–94, 1922.
- [19] I. Campbell, "Chi-squared and fisher-irwin tests of two-by-two tables with small sample recommendations," *Statistics in medicine*, vol. 26, no. 19, pp. 3661–3675, 2007.